

Fundamentals of computational biology

Lecture notes

Camilo García-Botero

2022-04-13T00:00:00-05:00

Table of contents

Preface	5
Introduction	6
I Unix	7
1 The command line	8
Getting started with the command line	8
1.1 Basic UNIX commands	8
1.2 Advance UNIX commands	8
1.3 Conda environments and other package managers	8
1.4 Notions of High Performance Computing (HPC)	8
II Sequence analysis and genomics	9
2 Introduction to sequence analysis	10
Endless debate: bioinformatics vs. computational biology	10
The duality of DNA	10
The central dogma theory of molecular biology extended	10
Sequencing strategies	10
Sequencing over time	10
Some insights from sequencing genomes	10
3 Sanger analysis	11
Databases exploration	11
Sanger sequencing methods	11
Files from Sanger	11
Sanger processing workflow	11
The 16S rRNA and its relevance for sequencing	11
4 Sequence alignments	12
Why do we align sequences?	12
What is homology	12
Pairwise alignments algoritms	12

The genetic code and Scoring matrices	12
BLAST and its families	12
Multiple sequence alignments	12
5 phylogenetic reconstructions	13
What is a phylogenetic tree	13
Methods for phylogenetic reconstruction	13
Building a phylogenetic reconstruction	13
6 NGS and TGS: principles	14
Platforms yields	14
Reads main differences	14
Illumina principle (sequencing by synthesis)	14
PacBio principle (sequencing by incorporation)	14
Oxford Nanopore Technology (ONT) principle	14
7 Genome assembly	15
The problem of assembling genomes	15
Main algorithms for genome assembly	15
Main concepts of an assembly	15
A complete workflow for assembling genomes	15
Assessing genomes	15
Understanding genome difficulties	15
8 Genome annotation and visualizaiton	16
8.1 <i>ab initio</i> annotation	16
8.2 Homolgy annotation	16
8.3 Annotation files	16
8.4 Visualizing genomes and annotations	16
9 Variant calling analysis	17
Common mutations	17
Structural variants	17
Genome rearrengments	17
Read mapping algorithms and programs	17
Identifying mutations	17
III Structural bioinformatics	18
10 Introduction to structural bioinformatics	19

IV Challenges	20
11 Genome searching	21
Challenge	21
12 Sanger processing	24
Challenge	24
Processing a single .ab1 pair	24
Processing a bulk of .ab1 files	25
13 Sequence alignment demo	26
Challenge	26
Download sequences	26
Unwrapping FASTA records	26
Gene search	27
Renaming fasta headers	27
Sequence alignment	28
Assesment of the alignment	28
An alternative approach using BLAST	29
The alternative using the GCF	29
14 Phylogenetic reconstruction	30
Challenge	30
Sequence alignment <i>cytb</i>	30
Evolutionary substitution model	30
Maximum likelihood reconstruction	30
Bayes inference reconstruction	31
15 Sequence reads assesment	32
Challenge	32
General stats from fastq files	32
A graphical assessment of reads	32
16 Genome assembly	33
Challenge	33
Download the reads	33
Assess read qualities	33
Exploring assemblers	33
References	35

Preface

We started this book with the aim of compiling the lectures of the course Fundamentals of Computational Biology offered at Universidad EAFIT for undergrad students in Biology. The course has been taught from different perspectives from its creation, yet the last iteration was divided into three modules. i) introduction to Unix (4 lectures) ii) introduction to sequence analysis and genomics (7 lectures) and iii) principles of structural biology (4 lectures).

Lectures are focused on a theoretical-practical approach where basic concepts from biology, bioinformatics and computer science and interleave with the practice to solve challenges.

Introduction

Here we present a course centered book of the Fundamentals of Computational Biology. We will cover several topics, from using the unix tools, the importance of package manager systems (such as homebrew and conda), sequencing technologies, sequence alignments, molecular phylogenetics, genome assembly and annotation, and variant calling analysis.

Part I

Unix

1 The command line

In this chapter we will explore the fundamentals of the command line. That is the concepts of Unix based systems the command line (CLI) and how we can use it to access information programmatically.

Getting started with the command line

1.1 Basic UNIX commands

1.2 Advance UNIX commands

1.3 Conda environments and other package managers

1.4 Notions of High Performance Computing (HPC)

Part II

Sequence analysis and genomics

2 Introduction to sequence analysis

In this chapter we will discuss several about several points of view about bioinformatics and computational biology and how to get started with the command line being a biologist, we will further consider several biological concepts that appear central to understand the manipulation of biological data.

Endless debate: bioinformatics vs. computational biology

The duality of DNA

The central dogma theory of molecular biology extended

Sequencing strategies

Sequencing over time

Some insights from sequencing genomes

3 Sanger analysis

This is a section about the first gen sequencing tech

Databases exploration

Sanger sequencing methods

The chain termination method

Sanger with capillary electrophoresis

Strengths and limitations of Sanger methods

Files from Sanger

Sanger processing workflow

The 16S rRNA and its relevance for sequencing

4 Sequence alignments

Why do we align sequences?

In search of homology and identity

What is homology

Pairwise alignments algorithms

Hamming distance

Edit distance

4.0.0.1 Dynamic programming

Needleman-Wunsch (global alignment)

Smith-Waterman (local alignment)

The genetic code and Scoring matrices

BLAST and its families

Multiple sequence alignments

5 phylogenetic reconstructions

What is a phylogenetic tree

Methods for phylogenetic reconstruction

Building a phylogenetic reconstruction

Evolutionary substitution model

Maximum likelihood

Bayesian inference

6 NGS and TGS: principles

Platforms yields

Reads main differences

Illumina principle (sequencing by synthesis)

The fastq format

Quality assesment of Illumina

PacBio principle (sequencing by incorporation)

Throughput evolution

Quality assesment of PacBio

Oxford Nanopore Technology (ONT) principle

Platforms

The fast5 file format

7 Genome assembly

The problem of assembling genomes

Main algorithms for genome assembly

Overlay, Layout, Consensus (OLS)

De Bruijn graphs

Main concepts of an assembly

Contigs, Unitigs, Scaffolds

A complete workflow for assembling genomes

Assessing genomes

Inspecting genome graphs

Genome completeness

Understanding genome difficulties

- End of chromosomes
- Errors
- Lack of coverage
- Heterozygosity
- repeats

8 Genome annotation and visualizaiton

8.1 *ab initio* annotation

8.2 Homolgy annotation

8.3 Annotation files

8.3.1 the GBK and GBFF

8.3.2 The GFF specificalitons

8.4 Visualizing genomes and annotations

9 Variant calling analysis

Common mutations

Structural variants

Genome rearrangements

Read mapping algorithms and programs

Burrow-Wheeler-Alignment

BWA-MEM2

Minimap2

SAM, BAM and CRAM formats

Identifying mutations

Freebayes and Snippy

The VCF file

Part III

Structural bioinformatics

10 Introduction to structural bioinformatics

Part IV

Challenges

11 Genome searching

Challenge

Your profesor is interested on knowing how many complete genomes of *Bacillus subtilis* are there in the NCBI databases. He ask you later to count the number of features (genes, CDS, ncRNA, rRNA, etc.) in the genome of *Bacillus subtilis* NCIB 3610 (GCF_002055965.1). And tell you to document each of the steps and how did you end up with the answer. Saving the file with your initials (e.g., CG-activity01.{md,txt,docx})

Downloading a genome

`ncbi-genome-download`

Downloading from NCBI

The first step in this journey is to download a bunch of sequences programatically. To do so, we will use the program `ncbi-genome-download`.

You could inspect all the options it provides, now we will set our command as the following:

```
#| echo: true
#| eval: true
ngd --genera "Bacillus subtilis" \
  -s refseq \
  -l complete \
  -o Data \
  --flat-output \
  --format features \
  -n bacteria \
  | head -n 10
```

Listing files

```
ls Data | head -n 10
```

Decompressing using gzip

```
gzip -d *
```

...

Some files in our data dir

```
ls Data | head
```

Importing the files into R

```
library(tidyverse)
library(fs)

all_features <- dir_ls("Data/") %>%
  map_df(read_tsv)

all_features %>%
  head()
```

...

```
library(tidyverse)
```

```
-- Attaching packages ----- tidyverse 1.3.1 --
```

```
v ggplot2 3.3.5      v purrr   0.3.4
v tibble  3.1.6      v dplyr   1.0.8
v tidyr   1.2.0      v stringr 1.4.0
v readr   2.1.2      v forcats 0.5.1
```

```
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()     masks stats::lag()
```

```
library(fs)

all_features <- dir_ls("Data/") %>%
  map_df(read_tsv)

all_features %>%
  head()
```

```
# A tibble: 0 x 0
```

Data processing

```
all_features_grouped <- all_features %>%
  rename(feature = `# feature`) %>%
  select(assembly, feature) %>%
  group_by(assembly, feature) %>% operations
  count() %>%
  pivot_wider(names_from = feature, values_from = n) %>%
  arrange(desc(CDS))

all_features_grouped %>%
  head()
```

create a new dataset that will group by features per accession. get read of the weird name of the column. Select these two columns. Group by these two columns to perform. count the numbers of rows based on the applied group. generate a wide dataset sending row names as columns. Arrange descending by the number of CDSs.

12 Sanger processing

Challenge

Your professor gives you a couple of .ab1 files of a 16S rRNA gene from an old project a student conducted. She tells you to process and analyse them using the Sanger sequence pipeline analysis. And as she doesn't know from which species they belong, she ask you to identify the organism to whom it belongs by using the resulting consensus sequence. She finally reminds you to document each step of the process including the identification step.

Processing a single .ab1 pair

```
library(sangeranalyseR)

groEL <- SangerAlignment(
  ABIF_Directory = "~/Projects/Bacillus/Data/Sanger/Inter/groEL/",
  REGEX_SuffixForward = "_1_F.ab1",
  REGEX_SuffixReverse = "_2_R.ab1",
  TrimmingMethod = "M2",
  M2CutoffQualityScore = 33,
  M2SlidingWindowSize = 10
)

writeFasta(groEL,
  outputDir = "~/Documents/Teaching/BiologyCourses/BI0487/Demos/02-demo-sangeranalysis",
  selection = "contigs_unalignment",
)

launchApp(groEL)
generateReport(groEL)
qualityBasePlot(groEL)
```


Processing a bulk of .ab1 files

```
library(fs)
library(purrr)

dirs <- fs::dir_ls("~/Projects/Bacillus/Data/Sanger/Inter")

sanger_bulk <- function(dir) {
  SangerAlignment(
    ABIF_Directory = dir,
    REGEX_SuffixForward = "_1_F.ab1",
    REGEX_SuffixReverse = "_2_R.ab1"
  )
}

genes <- dirs %>%
  map(sanger_bulk)

launchApp(genes$~/Users/camilogarcia/Projects/Bacillus/Data/Sanger/Inter/gyrA`)

writeFasta(
  outputDir = "~/Documents/Teaching/BiologyCourses/BI0487/Demos/02-demo-sangeranalysis",
  selection = "contigs_unalignment"
)
```

13 Sequence alignment demo

Challenge

Your professor is working with species from genus *Bacillus* and want to align an orthologous gene from 10 genomegits of different isolates. He gives you the GenBank accession number of these isolates and ask you to select one orthologous gene (Nucleotide seq) that you consider might be useful to differentiate the bacterial isolates and ask you to align those genes as you better consider. He finally ask you to document each step and send him the sequence alignment file in FASTA format along with the sequence alignment general stats in a TXT file (length, number of each nucleotides and other stats you consider important).

Accessions: GCA_012225885.1, GCA_000196735.1, GCA_000742895.1, GCA_001584335.1, GCA_000007825.1, GCA_000832905.1, GCA_000008425.1, GCA_000507105.1, GCA_000832605.1, GCA_900186955.1

Download sequences

Make sure to use the `--flat-output` avoiding download of multiple metadata

```
ngd --flat-output -p 4 -s genbank -A genome-accessions.txt -F cds-fasta bacteria
```

In this case `cds-fasta` parameter will download the nucleotide sequences of the gene. Other alternatives could be useful such as blast search on a genome database or searching through the GENBANK annotation files (both files also could be downloaded using `ngd`).

Unwrapping FASTA records

NCBI registries came with an undesirable wrapping around the lines of sequencing which basically is inserting a return character after some established number of characters. Then a way to get rid of them is to use a command line utility from [AstroBioMike \(Mike Lee\)](#) which will give a line per sequence after the FASTA header. We can later assume the the first line after the header will be the entire sequence

```
for i in GCA_*; do
    N=$(basename $i .fna);
    bash bit-remove-wraps.sh ${i} > ${N}_unwrapped.fasta;
done
```

Gene search

A possible way to search throughout the file registries is by using the **grep** command, that recursively will search each file. Fine tuned it allow to search for the first match, but also for the “after-context” in terms of lines desired to be printed:

```
grep -h\
-m 1\
-A 1\
-E "DNA gyrase, A| gyrase subunit A | gyrase alpha| gyrase \(\subunit A\)| gyrA" *.fasta
sed "s/--//g" | \
sed "/^$/d"
```

After finding the genes we could exclude some lines using **sed** avoid the “-” characters and the empty blank line using the appropriate regular expression (`^$/d`) . We are now with an almost clean multi sequence file, because header names are still and will be problematic. How do we programmatically change the FASTA headers? We will see in the next step.

Renaming fasta headers

A simple but powerful script to do this is **bit-dedup-fasta-heades** it was developed by [AstroBioMike \(Mike Lee\)](#) and it simply parses the headers and substitutes by a simple encoder found en each of them:

```
python bit-dedupe-fasta-headers.py -i all_gyrA.fasta -o all_gyrA_renamed.fasta
```

Now the the files has files names that are simply to work with. Which will enable to asses better out sequence alignment matrix.

Sequence alignment

There are many programs that are suited for performed multiple sequence alignments. Perhaps the two most used are [MAFFT](#) and [MUSCLE](#) both specialized in multiple sequence alignment (that is: when having two or more than two sequences). The second tends to be more accurate when having large data-sets, but the first on is more versatile, fast and accurate on different kind of data-sets.

Both program take as input a single file containing all the sequences concatenated horizontally (that is a multi-fasta file) careless of the extension but (MFA, FA, FASTA, FNA, etc). And generate a simple output (whether with the `-o` in [MUSCLE](#) or to the std output in [MAFFT](#))

```
ginsi --preservecase --reorder all_gyrA_renamed.fasta > all_gyrA_renamed_ginsi.fasta # global
einsi --preservecase --reorder all_gyrA_renamed.fasta > all_gyrA_renamed_einsi.fasta # gene-
linsi --preservecase --reorder all_gyrA_renamed.fasta > all_gyrA_renamed_linsi.fasta # local

muscle -i all_gyrA_renamed.fasta -o all_gyrA_renamed_muscle.fasta

famsa -t 8 all_gyrA_renamed.fasta > all_gyrA_renamed_famsa.fasta

kalign -i all_gyrA_renamed.fasta -o all_gyrA_renamed_kalign.fasta
```

Assesment of the alignment

Inspection of the alignment is there very first step for assesing its quality. A CDS tends to generate a codon-like alignment starting with the methione codon (ATG,GTG) and finishing with a stop (TAA, TAG, etc.). Therefore finding this structure when aligning a complete genes is expected. If a middle fraction of the gene is being aligned ORF might not display any stop codon. Verifying a codon-like alignment shows a biological order on the sequences other that mere artifact of the alignment, that is an evolutionary behavior of the sequence. We can do it usin [seqfu](#) from the CLI or interactively with [AliView](#).

A second step is to find the variability of the alignment. A simple way to find that is to calculate simpl stats from the alignment (sites, variable sites, As, Ts, etc.). A powerful cli program to do so is [goalign](#)

```
goalign stats -i all_gyrA_renamed_linsi.fasta
```

```
length 2508
nseqs 8
avgalleles 1.7400
variable sites 1202
char nb freq
- 273 0.013606
A 6418 0.319876
C 3633 0.181071
G 4755 0.236992
T 4985 0.248455
alphabet nucleotide
```

An alternative approach using BLAST

```
ngd --flat-output -p 4 -s genbank -A genome-accessions.txt -F fasta --parallel 8 bacteria
for i in GCA_*; do cat ${i} >> all_genomes.fasta; done
makeblastdb -in all_genomes.fasta -parse_seqids -blastdb_version 5 -title "demo" -dbtype nuc
blastn -db all_genomes.fasta -query gyrA.fasta -outfmt "6 sseqid sseq" -word_size 5 -evaluate 1
```

The alternative using the GCF

```
for i in *fna; do; goalign subset -e "gyrA" -i ${i} --unaligned;done | grep ">"
```

14 Phylogenetic reconstruction

Challenge

Your professor has been working with some mammal species and want to know the relationships of some sampled individuals. To do so he extracted the DNA and amplified the mitochondrial CYTB gene of those individuals. He gives a folder with multiple sequences and ask you to align them and to reconstruct two trees one using maximum likelihood (ML) and other using a Bayesian inference (BI). Then ask you to explain if both trees are congruent with each other.

Sequence alignment *cytb*

```
linsi --preserve-case --reorder cytb.fasta > cytb-aligned.fasta
```

Evolutionary substitution model

```
modeltest-ng -i cytb-aligned.fasta -d nt -o model-cytb.txt
```

```
04-demo-phylogenetics/model-cytb.txt.log
```

Maximum likelihood reconstruction

```
raxml-ng --msa cytb-aligned.fasta --model GTR+I+G4
```

Tree building

```

raxml_data <- read.tree("cytB-aligned.fasta.raxml.support")

raxml_data$tip.label <- str_replace_all(raxml_data$tip.label, "_", " ")

(
  raxml_tree <- ggtree(raxml_data) +
    geom_tiplab() +
    # geom_point2(aes(subset = !isTip, fill = as.integer(label)), shape = 24, size = 3) +
    geom_text(aes(label = as.integer(label))) +
    theme_tree(legend.position = c(0.8, 0.7))
)

```

Bayes inference reconstruction

Tree building

```
tail -n 18 ../Data/cytB-mb.nex
```

```
mb -i ../Data
```

```
sumtrees.py -s mcct -o=cytB-mb-mcct.tre cytB-aligned.fasta-out.nex.run1.t
```

15 Sequence reads assesment

Challenge

Your professor challenges you to assembly a bacterial genome. He wants to know if an assembly using Illumina reads or Nanopore reads is better. The raw sequences for Illumina could be downloaded with the code SRR15634574. And the Nanopore raw reads with SRR15634573 here: <https://sra-explorer.info>. Choose at least on set and follow the instruction from the lecture to assemble the genome. Document each step and send to the professor the `assembly-{your-initials}.fasta`. You could form groups of max. 4 students.

General stats from fastq files

A simple but fast cli to display the general stats from fastq is `seqfu`

```
seqfu stats -n *.gz
```

A graphical assessment of reads

Often its important to assess reads graphically and `nanoplot` offers a complete graphical summary (and general stats as well) of the reads

```
nanoplot --fastq reads.fastq.gz --output reads-report
```


16 Genome assembly

Challenge

Your professor has sequenced a bacterial isolate using PacBio and Nanopore sequencing methods and has got the FASTQ files from both technologies. Now he needs to know the quality and quantity of these data before start any other analysis and ask you to assess the data. He needs to know how many sequences there are, how many base pairs (in GB) are there and the N50. He is also interested in see a visualization of the i) number of bases vs. sequence lengths (log transformed) and ii) the read length vs. read quality vs. read number.

He ask you to document every step and to conclude what data should be used.

Download the reads

```
wget
```

Assess read qualities

When using illumina `fastqc` is a very fast alternative. For nanopore `nanplot` will do the job.

```
fastqc
```

Exploring assemblers

One of the most popular genome assemblers for NGS is `spades` whereas for TGS `flye` has been widely used

Shovill: spades under the hood

`shovill` is a pipeline that enables pre and post processing of genomic data. It can be tuned to several tools for the processing steps and also to select different standalone assemblers

```
shovill --outdir MxanthusIllumina\  
        --R1 R1.fq.gz\  
        --R2 R2.fq.gz\  
        --trim\  
        --cpus 32
```

Dragonflye: flye under the hood

Similar to `shovill` (and inspired by it) `dragonflye` is a pipeline that enables several processing steps of genomic data be

```
dragonflye --outdir MxanthusNanopore\  
           --gsize 9Mb\  
           --trim\  
           --reads ont-readsfastq.gz\  
           --racon 5
```

```
wget
```

Since we are trying to assemble a bacterial genome, computer memory appears to be a limiting features of a local machine. Then, a computer cluster with high performance turns out to be an important need.

First we need environment installations, therefore its important to have conda environments with the assemblers and other programs (`conda create -c bioconda dragonflye dragonflye` and `conda create -c bioconda shovill shovill`). That way both assemblers pipelines will lie in separate environment avoiding possible dependencies problems

We will use Apolo computer cluster which uses Slurm as the computer system workload manager (i.e a program that manages the time and resources of the computer).

References